

Artificial Intelligence Lab Work (4)
レポート解答用紙 (Report Answer Sheet)

学生証番号 (Student ID): 20521150

名前(Name): Pham Quoc Cuong

問題 1.

(プログラム)

Q1. Implement the MNIST learning and inference program by following the 10th lecture's slides (copy the program on the slide), and submit the program (.py) and the execution results, i.e., loss at each epoch during training and accuracy against test data, displayed on the console in a word file.

```
[55] import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn.functional as F
import torchvision as tv
import seaborn as sns
from sklearn.metrics import confusion_matrix

train_dataset=tv.datasets.MNIST(root=".",train=True,transform=tv.transforms.ToTensor(),download=True)
test_dataset=tv.datasets.MNIST(root=".",train=False,transform=tv.transforms.ToTensor(),download=True)

train_loader=torch.utils.data.DataLoader(dataset=train_dataset,batch_size=100,shuffle=True)
test_loader=torch.utils.data.DataLoader(dataset=test_dataset,batch_size=100,shuffle=False)
```

```
[56] MODELNAME = 'mnist.model'
EPOCH = 10
DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
[57] class MNIST(torch.nn.Module):
    losses = []

    def __init__(self):
        super(MNIST,self).__init__()
        self.l1 = torch.nn.Linear(784,300)
        self.l2 = torch.nn.Linear(300,300)
        self.l3 = torch.nn.Linear(300,10)

    def forward(self,x):
        h = F.relu(self.l1(x))
        h = F.relu(self.l2(h))
        y = self.l3(h)
        return y
```

```
[58] def train():
    model = MNIST().to(DEVICE)
    optimizer = torch.optim.Adam(model.parameters())
    for epoch in range(EPOCH):
        loss = 0
        for images, labels in train_loader:
            images = images.view(-1, 28*28).to(DEVICE)
            labels = labels.to(DEVICE)
            optimizer.zero_grad()
            y = model(images)
            batchloss = F.cross_entropy(y, labels)
            batchloss.backward()
            optimizer.step()
            loss = loss + batchloss.item()
        print('epoch: ', epoch, ', loss: ', loss)
        model.losses.append(loss)
    torch.save(model.state_dict(), MODELNAME)

    plt.plot(model.losses)
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Training Loss')
    plt.show()
```

```

[60] def test():
    total = len(test_loader.dataset)
    correct = 0
    model = MNIST().to(DEVICE)
    model.load_state_dict(torch.load(MODELNAME))
    model.eval()

    true_labels = []
    pred_labels = []

    for images, labels in test_loader:
        images = images.view(-1, 28 * 28).to(DEVICE)
        y = model(images)
        labels = labels.to(DEVICE)
        pred = y.max(dim=1)[1]

        true_labels.extend(labels.cpu().numpy())
        pred_labels.extend(pred.cpu().numpy())

    correct += (pred == labels).sum()

    print('Correct: ', correct.item())
    print('Total: ', total)
    print('Accuracy: ', correct.item() / float(total))

    # Generate confusion matrix
    cm = confusion_matrix(true_labels, pred_labels)


    # Plot the confusion matrix as an image
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

```

(実行結果)


```
[59] train()
```

```
epoch: 0 , loss: 171.65008279308677
epoch: 1 , loss: 65.32018205523491
epoch: 2 , loss: 42.61820778762922
epoch: 3 , loss: 31.041309498250484
epoch: 4 , loss: 23.00104690855369
epoch: 5 , loss: 18.023108201567084
epoch: 6 , loss: 14.666266271553468
epoch: 7 , loss: 12.384661977353971
epoch: 8 , loss: 9.112057573875063
epoch: 9 , loss: 9.07755283745064
```



```
[61] test()
```

```
Correct: 9786
Total: 10000
Accuracy: 0.9786
```



問題 2

(プログラム)

Q2. Rewrite the program you wrote in Q1 to train and infer on the image recognition dataset CIFAR-10, and submit the program (.py) and the execution results, i.e., loss at each epoch during training and accuracy against test data, displayed on the console in a word file. CIFAR-10 is a 10-class image classification data, and can be downloaded by the following program.

```
[62] train_dataset = tv.datasets.CIFAR10(root=".", train=True, transform=tv.transforms.ToTensor(), download=True)
    test_dataset = tv.datasets.CIFAR10(root=".", train=False, transform=tv.transforms.ToTensor(), download=True)

    train_loader=torch.utils.data.DataLoader(dataset=train_dataset,batch_size=100,shuffle=True)
    test_loader=torch.utils.data.DataLoader(dataset=test_dataset,batch_size=100,shuffle=False)
```

```
Files already downloaded and verified
Files already downloaded and verified
```

```
[63] MODELNAME = 'cifar-10.model'
    EPOCH = 10
    DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
[64] class CIFAR(torch.nn.Module):
    losses = [] # List to store the loss values

    def __init__(self):
        super(CIFAR,self).__init__()
        self.l1 = torch.nn.Linear(32*32*3,300)
        self.l2 = torch.nn.Linear(300,300)
        self.l3 = torch.nn.Linear(300,10)

    def forward(self,x):
        h = F.relu(self.l1(x))
        h = F.relu(self.l2(h))
        y = self.l3(h)
        return y
```

```
[65] def train():
    model = CIFAR().to(DEVICE)
    optimizer = torch.optim.Adam(model.parameters())
    for epoch in range(EPOCH):
        loss =0
        for images, labels in train_loader:
            images = images.view(-1,32*32*3).to(DEVICE)
            labels = labels.to(DEVICE)
            optimizer.zero_grad()
            y = model(images)
            batchloss = F.cross_entropy(y, labels)
            batchloss.backward()
            optimizer.step()
            loss = loss + batchloss.item()
        print('epoch: ', epoch, ', loss: ',loss)
        model.losses.append(loss)
    torch.save(model.state_dict(), MODELNAME)
    # Plot the losses result
    plt.plot(model.losses)
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Training Loss')
    plt.show()
```

```

[67] def test():
    total = len(test_loader.dataset)
    correct = 0
    model = CIFAR().to(DEVICE)
    model.load_state_dict(torch.load(MODELNAME))
    model.eval()

    true_labels = []
    pred_labels = []

    for images, labels in test_loader:
        images = images.view(-1, 32 * 32 * 3).to(DEVICE)
        y = model(images)
        labels = labels.to(DEVICE)
        pred = y.max(dim=1)[1]

        true_labels.extend(labels.cpu().numpy())
        pred_labels.extend(pred.cpu().numpy())

    correct += (pred == labels).sum()

    print('Correct: ', correct.item())
    print('Total: ', total)
    print('Accuracy: ', correct.item() / float(total))

    # Generate confusion matrix
    cm = confusion_matrix(true_labels, pred_labels)


    # Plot the confusion matrix as an image
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

```

(実行結果)

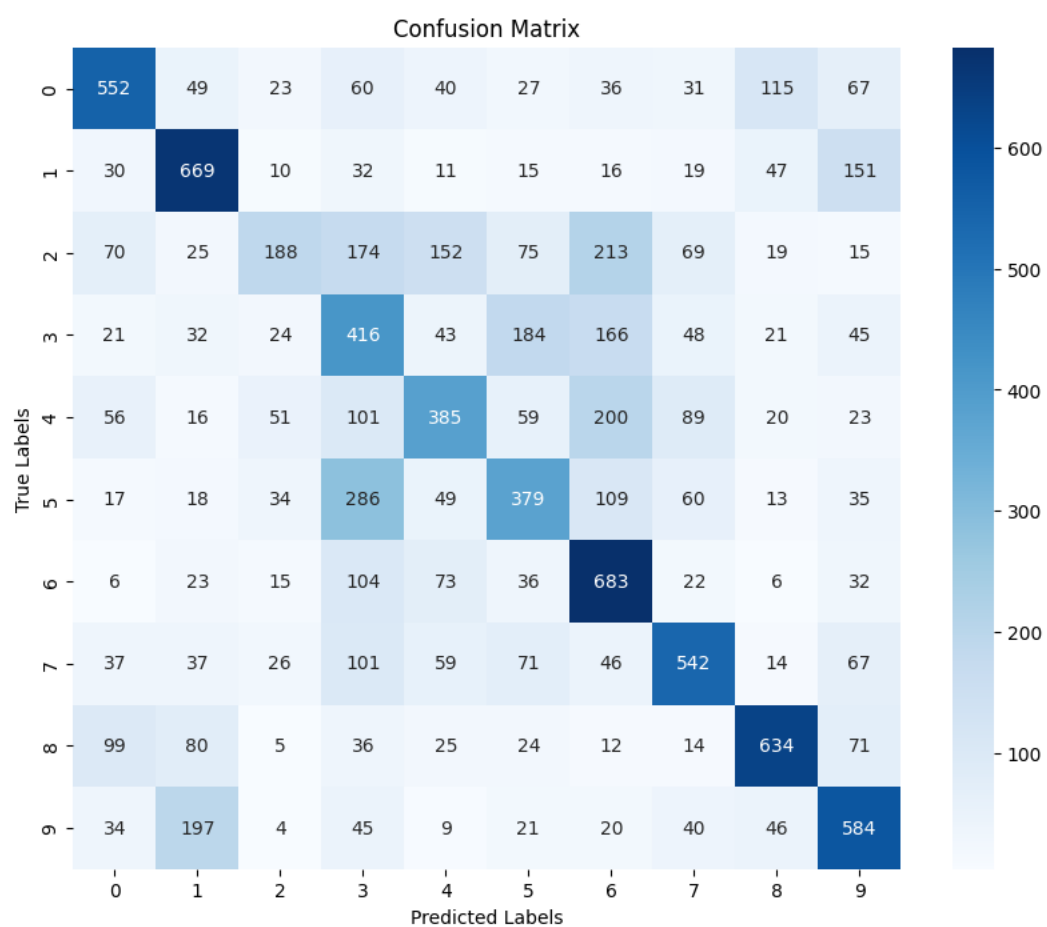
```
[66] train()
```

```
epoch: 0 , loss: 924.9251463413239
epoch: 1 , loss: 831.4743949174881
epoch: 2 , loss: 794.2123394012451
epoch: 3 , loss: 759.8124611377716
epoch: 4 , loss: 741.5656607151031
epoch: 5 , loss: 719.1269598007202
epoch: 6 , loss: 707.9702979326248
epoch: 7 , loss: 691.9842377901077
epoch: 8 , loss: 680.1186082363129
epoch: 9 , loss: 667.5125402212143
```



```
[68] test()
```

```
Correct: 5032
Total: 10000
Accuracy: 0.5032
```

問題 3

(プログラム)

Q3. Rewrite the program you wrote in Q2 to create an NN with one intermediate layer as a convolutional layer, and submit the program (.py) and the execution results, i.e., loss at each epoch during training and accuracy against test data, displayed on the console in a word file. The convolutional layer can be obtained by the following program

nn.Conv2d(in_channel, out_channel, filtersize)

where in_channel is the number of input channels, out_channel is the number of output channels, and filtersize is the size of the filter.

```
[69] MODELNAME = 'cifar-10-conv2d.model'
    EPOCH = 10
    DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'

[70] class CIFAR_Conv2D(torch.nn.Module):
    losses = [] # List to store the loss values

    def __init__(self):
        super(CIFAR_Conv2D, self).__init__()
        self.l1 = torch.nn.Conv2d(3, 100, 4)
        self.l2 = torch.nn.Linear(29*29*100, 300)
        self.l3 = torch.nn.Linear(300, 10)

    def forward(self, x):
        h = F.relu(self.l1(x))
        h = torch.flatten(h, start_dim=1)
        h = F.relu(self.l2(h))
        y = self.l3(h)
        return y

[71] def train():
    model = CIFAR_Conv2D().to(DEVICE)
    optimizer = torch.optim.Adam(model.parameters())

    for epoch in range(EPOCH):
        loss = 0
        for images, labels in train_loader:
            images = images.view(-1, 3, 32, 32).to(DEVICE)
            labels = labels.to(DEVICE)
            optimizer.zero_grad()
            y = model(images)
            batchloss = F.cross_entropy(y, labels)
            batchloss.backward()
            optimizer.step()
            loss = loss + batchloss.item()

        print('epoch: ', epoch, ', loss: ', loss)
        model.losses.append(loss) # Append the loss to the list

    torch.save(model.state_dict(), MODELNAME)

    plt.plot(model.losses)
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Training Loss')
    plt.show()
```

```
[73] def test():
    total = len(test_loader.dataset)
    correct = 0
    model = CIFAR_Conv2D().to(DEVICE)
    model.load_state_dict(torch.load(MODELNAME))
    model.eval()

    true_labels = []
    pred_labels = []

    for images, labels in test_loader:
        images = images.view(-1, 3, 32, 32).to(DEVICE)
        y = model(images)
        labels = labels.to(DEVICE)
        pred = y.max(dim=1)[1]

        true_labels.extend(labels.cpu().numpy())
        pred_labels.extend(pred.cpu().numpy())

    correct += (pred == labels).sum()

    print('Correct: ', correct.item())
    print('Total: ', total)
    print('Accuracy: ', correct.item() / float(total))


    # Generate confusion matrix
    cm = confusion_matrix(true_labels, pred_labels, labels=range(10))

    # Plot the confusion matrix as an image
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=range(10), yticklabels=range(10))
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()
```

(実行結果)

```
[72] train()
```

```
epoch: 0 , loss: 812.6806900501251
epoch: 1 , loss: 570.0182233452797
epoch: 2 , loss: 483.50918436050415
epoch: 3 , loss: 410.78025180101395
epoch: 4 , loss: 339.84813117980957
epoch: 5 , loss: 267.78175631165504
epoch: 6 , loss: 201.96683129668236
epoch: 7 , loss: 145.0935261696577
epoch: 8 , loss: 101.98463297635317
epoch: 9 , loss: 75.8632026873529
```



```
[74] test()
```

```
Correct: 6178
Total: 10000
Accuracy: 0.6178
```

